

Week 3: three rounds of choosing

Emir Cabalak · General AI Fluency · Consistency, Not Talent (and Frame, Not Upstage)

The week-3 package in one document: the one-line claim, the content map with its gather list, the identity kit, and the image set. I have arranged it around what the session was about rather than around the three cards, because the same move happens three times. Generate many, keep one, be able to say why the others died. Roughly twenty options were produced across the week and four survived. What follows is mostly the rejections, since that is where the judgment is visible.

ROUND ONE, THE WORDS

Ten claims generated, one kept.

I build machine learning models and then try to break my own results before anyone else can.

The cuts, and the rule each one taught me.

Rigorous, reproducible ML evaluation for data teams. Dead on arrival. That is the vocabulary of somebody who has never had to defend a number in a meeting, and it would sit unchanged on any consultancy page.
Turning messy search data into models that do not lie to you. Dies on "messy", which describes everyone's data and therefore describes nobody's.
Machine learning you can check. Also honest evaluation for models that have to survive production. True and forgettable. Both state a category rather than a behaviour, and nobody repeats a category.
I care more about the split than the algorithm. The sharpest line on the list and I nearly took it. Cut because it is an opinion about the field, not a claim about me. Good to say once somebody is already listening, bad as a first sentence.
ML engineer who reports the lower number. Closest runner up. Cut because "the lower number" needs a paragraph of setup, and the first sentence does not get one.
I find the leak in the pipeline, usually my own. The most human line here, and it undersells. It makes me sound like an auditor of other people's work. I build the model too and the claim has to say so.
The keeper. It names both halves of the job, building and breaking. "Break my own results" is specific and slightly uncomfortable, which is what makes it stick, and "before anyone else can" tells a reader what problem it solves for them. It could not be pasted onto somebody else's portfolio without becoming false, which is the only test that matters.

One sharpening pass after that. "ML" became "machine learning", because the first sentence on a page should not need an acronym. "Someone else does" became "anyone else can". And "try to" stayed, even though "and break" is punchier, because I do not always succeed and the whole page argues against overstating.

THE MAP THE CLAIM HAS TO SURVIVE

Five pages, one action, every arrow pointing at it.

The one action from week 1 is **read one case study all the way to the end**. So most calls to action point inward rather than at my inbox, and only the last one in the chain asks for anything.

Page	Sections, in order	Call to action
Home	name and claim, hero, the lead case, two more cases, one line about me, footer	Read the leakage case. It is the one where my score went down and I kept the lower number.
/leak leads	title, the problem, what I did, the decision, the population caveat, what came of it, screenshot	Case 2 is the other half of this.
/split	same shape, with why I published all three numbers and how history depth shrank the set	Case 3 is the same discipline applied to prompting.
/ladder	same shape, section 4 is the version that made things worse	If you want the code behind any number here, email me.
/about	portrait and bio, how I work, other work, tools, contact	Start with the leakage case.

Case 1 leads because it is the one that proves the claim directly: it is the case where I broke my own result. Case 3 goes last on purpose, because it is the least load bearing for somebody deciding on a hire. Nothing on the site competes with the one action. No newsletter, no contact form, no social row.

Still to gather, honestly. Three case screenshots, one per case page. Until they land the pages carry their numbers as real HTML tables, which is more accessible anyway. The held-out future month for case 2 does not exist yet either, and that case says so in its own limitations section rather than leaving a reader to assume otherwise.

ROUND TWO, THE FRAME

The design frames the work and never upstages it. Here is what that cost.

Headings: Space Grotesk


Weights 500 and 600. Free, Google Fonts.

I break my own results


Body: IBM Plex Sans
Weights 400 and 600. Drawn for technical documentation, so it survives a page full of digits.

ROC AUC dropped from 0.91 to 0.84 once the leaking pair came out, and 0.84 is the number I report.

What I cut. A third face, IBM Plex Mono, for the metric tables. It would have looked right. Three fonts is a pile, and font-variant-numeric: tabular-nums on Plex Sans aligns the digits without another download, which covers the actual reason I wanted mono.

			
Paper #FAF9F7 · background	Ink #16181D · all text · 17:1	Main #2E4B6B · rules, headers · 8.8:1	Accent #C2410C · the one CTA · 4.9:1

Contrast is measured against paper. All three pass WCAG AA for body text. The accent has exactly one job. If it starts appearing on subheads it stops meaning "click this" and the page loses its only pointer.

			
Paper #14161A	Ink #E9E7E3	Main #8FB2D6	Accent #F0834D
Dark mode. Same four roles lifted to survive a dark ground, behind a prefers-color-scheme block in tokens.css.			



Emir Cabalak

numbers that hold up



favicon.svg at 64, 32 and 16 pixels

logo.svg

The mark I rejected. A shield, then a checkmark. "Numbers you can trust" pushes you straight at trust badges, and both attempts looked like a security product. Claiming trustworthiness with an icon is the opposite of demonstrating it with a case study, so the monogram stayed plain and lowercase. EC in a box reads as a corporate badge; this is a person's site.


Fonts: Space Grotesk for headings, IBM Plex Sans for body. Colors: paper #FAF9F7, ink #16181D, main #2E4B6B, accent #C2410C, dark mode variants in tokens.css. Accent is reserved for the single call to action.

Mood: a lab notebook, not a brochure. Plain, quiet, and slightly technical, so the numbers and the writing are the loudest thing on the page. Generous white space, one column, no gradients, no shadows, no animation.


Why quiet is not a taste preference here. My claim is that my numbers hold up when somebody checks them. A page with gradients and motion undercuts that before a reader reaches a sentence, because it is visibly trying to impress and the whole argument is that I am not. That is the frame-not-upstage rule with something actually at stake: the design is making the same claim as the text, or it is arguing against it.

ROUND THREE, THE IMAGES


Nine slots. Three generated, six real or vector, five ideas killed.



hero-texture.svg. Somebody who knows what a train and test split looks like reads it in a second. Somebody who does not sees a quiet chart and reads the words.



icon-leak.svg
Accent on the leak.



icon-split.svg
Accent on the held-out side.



icon-ladder.svg
Accent on the rung that went down.

How the style is held steady. Not by describing a mood. By fixing the numbers: paper background, #2E4B6B strokes at two pixels, exactly one element in #C2410C and never more than one, no fills except the dots, no gradients, square line caps. Attempt one described the mood in words and generated all three at once, and produced three stroke weights and two different backgrounds. Describing a mood does not hold a style. Fixing the numbers does.

Where a real capture beats anything generated. The three case screenshots are the proof. A generated picture of a table of numbers is a picture of numbers I did not compute, on a site whose entire claim is that my numbers hold up. That is not a style preference, it is the argument. Rules for each: crop to the table, no browser chrome, no file paths or client names, light theme, 2x density so the digits survive.

The portrait is a real photograph. Plain wall, daylight, one frame, cropped to a circle. An AI portrait on a page arguing for honest evaluation is a joke that writes itself, and people can tell.

Slot	Page	Source
Hero texture	home	generated
Portrait	about	real photo, shot and live
Three case icons	home and each case	generated, one set
Feature verdict table	leak	real capture, to take
Three-split table	split	real capture, to take
V3 vs V4 outputs	ladder	real capture, to take
Favicon	every page	vector

WHAT I REJECTED, AND WHY EACH ONE DIED

The part that is actually graded.

A generated "data scientist at work" hero. Person at a desk, three monitors, blue code glow, the whole stock-photo grammar. It looked competent and it was the easiest thing on the page to delete. It says nothing my one-line claim does not say better, it makes a specific claim about a workspace that is not mine, and it sat in the most valuable space on the site doing decoration.
A generated neural network mesh. Glowing nodes and connecting lines. Rejected for a better reason than taste: it misdescribes the cases the hero sits above. I do build neural networks elsewhere and the about page links to that work, but these three cases are gradient boosting and cross-validation design. A hero makes a promise about the page underneath it, and the reader who came for deep learning leaves by paragraph two. An image that misdescribes the work is worse than no image.
A generated portrait. Killed on sight, for the reason above.
A shield, and then a checkmark, in the logo. Looked like a security product both times.
A third typeface. Wanted, justified, and still cut.
The pattern. Four of the five died for the same reason, and it is not "I liked the other one better". Each made a promise the page underneath could not keep. That is the only rejection test I ended up needing, and it works on words, on colour and on images equally.

FILES

Live at emircabalak.github.io. Identity kit: logo.svg, favicon.svg, tokens.css, which is the palette and the font stacks as CSS custom properties imported by the live site. Images: hero-texture.svg, icon-leak.svg, icon-split.svg, icon-ladder.svg, portrait.jpg. Written versions: identity-kit.md, image-set.md, content-map.md.